

# Mushroom Binary Classification

Name: Aniket Nitin Salunkhe

UCID: as4593

Dataset link:

[https://www.kaggle.com/datasets/prishasawhney/mushroom-dataset?select=mushroom\\_cleaned.csv](https://www.kaggle.com/datasets/prishasawhney/mushroom-dataset?select=mushroom_cleaned.csv)

GitHub Link: <https://github.com/AniketSalunkheNJIT/salunkhe-aniket-data-mining-finalProj>

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
import warnings
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split, StratifiedKFold,
KFold
from sklearn.metrics import confusion_matrix, precision_score,
recall_score, f1_score, roc_auc_score, accuracy_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Bidirectional,
Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
import time
warnings.filterwarnings("ignore")

# Load the dataset
file_path = 'mushroom_cleaned.csv'
data = pd.read_csv(file_path)

# Display dataset information
print("First 5 rows of the dataset:")
display(data.head())
print("\nDataset Info:")
data.info()
```

First 5 rows of the dataset:

	cap-diameter	cap-shape	gill-attachment	gill-color	stem-height \
0	1372	2	2	10	3.807467
1	1461	2	2	10	3.807467
2	1371	2	2	10	3.612496
3	1261	6	2	10	3.787572
4	1305	6	2	10	3.711971

	stem-width	stem-color	season	class
0	1545	11	1.804273	1
1	1557	11	1.804273	1
2	1566	11	1.804273	1
3	1566	11	1.804273	1
4	1464	11	0.943195	1

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 54035 entries, 0 to 54034

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	cap-diameter	54035 non-null	int64
1	cap-shape	54035 non-null	int64
2	gill-attachment	54035 non-null	int64
3	gill-color	54035 non-null	int64
4	stem-height	54035 non-null	float64
5	stem-width	54035 non-null	int64
6	stem-color	54035 non-null	int64
7	season	54035 non-null	float64
8	class	54035 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 3.7 MB

```
# Check for missing values
print("\nMissing Values in Dataset:")
print(data.isnull().sum())
```

Missing Values in Dataset:

cap-diameter	0
cap-shape	0
gill-attachment	0
gill-color	0

```
stem-height      0
stem-width       0
stem-color       0
season           0
class            0
dtype: int64
```


```
# Dataset statistics
print("\nDataset Statistics:")
display(data.describe())
```

Dataset Statistics:

	cap-diameter	cap-shape	gill-attachment	gill-color \
count	54035.000000	54035.000000	54035.000000	54035.000000
mean	567.257204	4.000315	2.142056	7.329509
std	359.883763	2.160505	2.228821	3.200266
min	0.000000	0.000000	0.000000	0.000000
25%	289.000000	2.000000	0.000000	5.000000
50%	525.000000	5.000000	1.000000	8.000000
75%	781.000000	6.000000	4.000000	10.000000
max	1891.000000	6.000000	6.000000	11.000000

	stem-height	stem-width	stem-color	season
count	54035.000000	54035.000000	54035.000000	54035.000000
mean	0.759110	1051.081299	8.418062	0.952163
std	0.650969	782.056076	3.262078	0.305594
min	0.000426	0.000000	0.000000	0.027372
25%	0.270997	421.000000	6.000000	0.888450
50%	0.593295	923.000000	11.000000	0.943195
75%	1.054858	1523.000000	11.000000	0.943195
max	3.835320	3569.000000	12.000000	1.804273

```
# Class distribution
sns.countplot(x='class', data=data, palette='coolwarm')
plt.title("Class Distribution")
plt.show()
```



```
# Separate features and target
X = data.drop(columns=['class'])
y = data['class']

# Standardize the dataset
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)
print("\nTraining and Testing Data Shapes:")
print(f"X_train: {X_train.shape} | y_train: {y_train.shape}")
print(f"X_test: {X_test.shape} | y_test: {y_test.shape}")
```

```
Training and Testing Data Shapes:
X_train: (43228, 8) | y_train: (43228,)
X_test: (10807, 8) | y_test: (10807,)
```

# Evaluation Metrics Description

To evaluate the performance of the models in this project, the following metrics are calculated:

1. **True Positives (TP):** The number of positive samples correctly classified as positive.
2. **True Negatives (TN):** The number of negative samples correctly classified as negative.
3. **False Positives (FP):** The number of negative samples incorrectly classified as positive.
4. **False Negatives (FN):** The number of positive samples incorrectly classified as negative.
5. **Total Positives (P):** The total number of actual positive samples in the dataset.
6. **Total Negatives (N):** The total number of actual negative samples in the dataset.

## Derived Metrics

1. **True Positive Rate (TPR):** Also called recall or sensitivity, it is calculated as  
$$\text{TPR} = \frac{\text{TP}}{\text{P}}$$
  
It measures the proportion of actual positives correctly identified.
2. **True Negative Rate (TNR):** Also called specificity, it is calculated as  
$$\text{TNR} = \frac{\text{TN}}{\text{N}}$$
  
It measures the proportion of actual negatives correctly identified.
3. **False Positive Rate (FPR):** The proportion of actual negatives incorrectly classified as positives, calculated as  
$$\text{FPR} = \frac{\text{FP}}{\text{N}}$$
4. **False Negative Rate (FNR):** The proportion of actual positives incorrectly classified as negatives, calculated as  
$$\text{FNR} = \frac{\text{FN}}{\text{P}}$$
5. **Precision:** The ratio of true positives to all predicted positives, calculated as  
$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$
6. **F1 Score:** The harmonic mean of precision and recall, calculated as  
$$\text{F1} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$
7. **Accuracy:** The overall proportion of correct predictions, calculated as  
$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{P} + \text{N}}$$
8. **Error Rate:** The proportion of incorrect predictions, calculated as  
$$\text{Error Rate} = 1 - \text{Accuracy}$$
9. **Balanced Accuracy:** The average of TPR and TNR, calculated as  
$$\text{Balanced Accuracy} = \frac{\text{TPR} + \text{TNR}}{2}$$
10. **True Skill Statistic (TSS):** The difference between TPR and FPR, calculated as  
$$\text{TSS} = \text{TPR} - \text{FPR}$$
  
It evaluates the skill of the classifier independent of class imbalance.

11. **Heidke Skill Score (HSS):** A metric assessing classifier accuracy while considering random chance.
12. **Brier Score:** Measures the mean squared error between predicted probabilities and the actual outcomes. Lower scores indicate better calibration.
13. **Brier Skill Score:** A normalized version of the Brier Score comparing the model to a baseline model (e.g., random guessing).

These metrics provide a comprehensive assessment of the model's performance, covering aspects such as precision, recall, and calibration.

```
from sklearn.metrics import confusion_matrix, roc_auc_score
import numpy as np

# Function to calculate metrics
def calculate_metrics(y_true, y_pred, y_prob=None):
    cm = confusion_matrix(y_true, y_pred)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]

    # Basic metrics
    P = TP + FN
    N = TN + FP
    TPR = TP / P if P > 0 else 0
    TNR = TN / N if N > 0 else 0
    FPR = FP / N if N > 0 else 0
    FNR = FN / P if P > 0 else 0
    precision = TP / (TP + FP) if (TP + FP) > 0 else 0
    f1 = 2 * precision * TPR / (precision + TPR) if (precision + TPR) > 0 else 0
    accuracy = (TP + TN) / (P + N) if (P + N) > 0 else 0
    error_rate = 1 - accuracy

    # Advanced metrics
    balanced_accuracy = (TPR + TNR) / 2
    tss = TPR + TNR - 1
    hss = 2 * (TP * TN - FP * FN) / ((P * (FP + TN)) + (N * (TP + FN))) if (P * (FP + TN) + N * (TP + FN)) > 0 else 0

    # Brier Score (requires probabilities)
    brier_score = np.mean((y_prob - y_true)**2) if y_prob is not None else None
    brier_skill_score = 1 - brier_score / np.var(y_true) if brier_score is not None else None

    # ROC AUC (requires probabilities)
    roc_auc = roc_auc_score(y_true, y_prob) if y_prob is not None else
```

None

```
return {
    'TP': TP, 'TN': TN, 'FP': FP, 'FN': FN,
    'P': P, 'N': N,
    'TPR': TPR, 'TNR': TNR, 'FPR': FPR, 'FNR': FNR,
    'Precision': precision, 'F1': f1, 'Accuracy': accuracy, 'Error
Rate': error_rate,
    'Balanced Accuracy': balanced_accuracy, 'TSS': tss, 'HSS':
hss,
    'Brier Score': brier_score, 'Brier Skill Score':
brier_skill_score,
    'ROC AUC': roc_auc # Added ROC AUC metric
}
```

## Random Forest Classification with Stratified K-Fold Cross-Validation

In this cell, we implement a Random Forest classification model and evaluate its performance using **Stratified K-Fold Cross-Validation**. This ensures that each fold maintains the same proportion of positive and negative samples, providing a more robust evaluation of the model's performance.

### Process

1. **Data Preparation:**
  - Convert `X_train` and `y_train` into NumPy arrays to enable indexing and slicing for Stratified K-Fold.
2. **Stratified K-Fold Cross-Validation:**
  - The dataset is split into 10 folds using `StratifiedKFold` with shuffling for randomness and a fixed random state for reproducibility.
  - For each fold:
    - The training and validation subsets are created.
    - A Random Forest model with 100 estimators is trained on the training subset.
    - Predictions (`y_pred_k`) and probabilities (`y_prob_k`) are generated for the validation subset.
3. **Metric Calculation:**
  - Custom evaluation metrics are calculated for each fold using the `calculate_metrics` function, which computes metrics like accuracy, precision, recall, F1-score, Brier score, and more.
4. **Metrics Aggregation:**
  - Metrics for each fold are stored in a list and later converted into a DataFrame for analysis.

- The mean of all folds is computed to derive the overall performance of the Random Forest model.

## Output

- **Fold-wise Metrics:** A detailed breakdown of the evaluation metrics for each fold.
- **Average Metrics:** The mean values of the metrics across all folds, representing the overall performance of the Random Forest model.

The results highlight the model's ability to handle the classification task, showing metrics such as **accuracy, precision, recall, F1-score, and Brier scores**, which provide a comprehensive view of its strengths and areas for improvement.

```
# Random Forest Model
rf_metrics = []
kf = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

for train_idx, val_idx in kf.split(X_train, y_train):
    X_train_k, X_val_k = X_train[train_idx], X_train[val_idx]
    y_train_k, y_val_k = y_train.iloc[train_idx],
y_train.iloc[val_idx]

    rf_model = RandomForestClassifier(random_state=42)
    rf_model.fit(X_train_k, y_train_k)
    y_pred_k = rf_model.predict(X_val_k)
    y_prob_k = rf_model.predict_proba(X_val_k)[: , 1]

    metrics = calculate_metrics(y_val_k, y_pred_k, y_prob_k)
    rf_metrics.append(metrics)

rf_metrics_df = pd.DataFrame(rf_metrics)
print("\nRandom Forest Metrics Per Fold:")
display(rf_metrics_df)

# Average Metrics for Random Forest
rf_avg_metrics = rf_metrics_df.mean()
print("\nRandom Forest Average Metrics:")
print(rf_avg_metrics)
```

## Random Forest Metrics Per Fold:

	TP	TN	FP	FN	P	N	TPR	TNR	FPR
FNR \									
0	2351	1925	24	23	2374	1949	0.990312	0.987686	0.012314
0.009688									
1	2352	1935	14	22	2374	1949	0.990733	0.992817	0.007183
0.009267									
2	2353	1924	25	21	2374	1949	0.991154	0.987173	0.012827
0.008846									
3	2354	1925	24	20	2374	1949	0.991575	0.987686	0.012314



0.008425									
4	2345	1924	25	29	2374	1949	0.987784	0.987173	0.012827
0.012216									
5	2343	1931	18	31	2374	1949	0.986942	0.990764	0.009236
0.013058									
6	2348	1927	22	26	2374	1949	0.989048	0.988712	0.011288
0.010952									
7	2350	1930	19	24	2374	1949	0.989890	0.990251	0.009749
0.010110									
8	2357	1931	17	17	2374	1948	0.992839	0.991273	0.008727
0.007161									
9	2355	1926	22	19	2374	1948	0.991997	0.988706	0.011294
0.008003									

	Precision	F1	Accuracy	Error Rate	Balanced Accuracy
TSS \					
0	0.989895	0.990103	0.989128	0.010872	0.988999
0.977998					
1	0.994083	0.992405	0.991672	0.008328	0.991775
0.983550					
2	0.989487	0.990320	0.989359	0.010641	0.989164
0.978327					
3	0.989907	0.990741	0.989822	0.010178	0.989631
0.979261					
4	0.989451	0.988617	0.987509	0.012491	0.987479
0.974957					
5	0.992376	0.989652	0.988665	0.011335	0.988853
0.977706					
6	0.990717	0.989882	0.988897	0.011103	0.988880
0.977760					
7	0.991980	0.990934	0.990053	0.009947	0.990071
0.980142					
8	0.992839	0.992839	0.992133	0.007867	0.992056
0.984112					
9	0.990745	0.991370	0.990514	0.009486	0.990351
0.980703					

	HSS	Brier Score	Brier Skill Score	ROC AUC
0	0.977998	0.010200	0.958803	0.999480
1	0.983550	0.008712	0.964814	0.999620
2	0.978327	0.009757	0.960590	0.999516
3	0.979261	0.010139	0.959048	0.999159
4	0.974957	0.010940	0.955813	0.998922
5	0.977706	0.011084	0.955233	0.998800
6	0.977760	0.010409	0.957960	0.999214
7	0.980142	0.010718	0.956709	0.999359
8	0.984112	0.009561	0.961382	0.999157
9	0.980703	0.009352	0.962224	0.999436

```
Random Forest Average Metrics:
TP          2350.800000
TN          1927.800000
FP           21.000000
FN           23.200000
P           2374.000000
N           1948.800000
TPR          0.990227
TNR          0.989224
FPR          0.010776
FNR          0.009773
Precision    0.991148
F1           0.990686
Accuracy     0.989775
Error Rate   0.010225
Balanced Accuracy 0.989726
TSS          0.979452
HSS          0.979452
Brier Score  0.010087
Brier Skill Score 0.959258
ROC AUC      0.999267
dtype: float64
```

## Naive Bayes Classification with Stratified K-Fold Cross-Validation

In this cell, we implement a **Naive Bayes classification model** and evaluate its performance using **Stratified K-Fold Cross-Validation**. This approach ensures consistent evaluation by maintaining the class distribution in each fold.

### Process

1. **Data Preparation:**
  - Convert `X_train` and `y_train` into NumPy arrays for easy manipulation and indexing during cross-validation.
2. **Stratified K-Fold Cross-Validation:**
  - The dataset is split into 10 folds using `StratifiedKFold`, ensuring the class balance is preserved in each fold.
  - For each fold:
    - The training and validation subsets are created.
    - A **Multinomial Naive Bayes (MultinomialNB)** model is trained on the training subset.
    - Predictions (`y_pred_k`) and probability scores (`y_prob_k`) are generated for the validation subset.
3. **Metric Calculation:**

- Custom metrics, such as accuracy, precision, recall, F1-score, Brier score, and more, are calculated for each fold using the `calculate_metrics` function.
  - Metrics for each fold are stored in a list.
4. **Metrics Aggregation:**
- Metrics for all folds are compiled into a DataFrame for detailed analysis.
  - The average metrics across all folds are computed to summarize the overall performance of the Naive Bayes model.

## Output

- **Fold-wise Metrics:** Displays the performance metrics for each fold, providing insights into the model's consistency across different splits.
- **Average Metrics:** Provides the mean values of the evaluation metrics across all folds, representing the overall capability of the Naive Bayes classifier.

## Observations

The results include key metrics such as:

- **Accuracy, Precision, Recall, F1-score** to measure classification performance.
- **Brier Score and Brier Skill Score** to evaluate the model's probability calibration.

These metrics provide a comprehensive assessment of the Naive Bayes classifier's performance, highlighting its strengths and limitations in the sentiment analysis task.

```
# Naive Bayes Model
nb_metrics = []
for train_idx, val_idx in kf.split(X_train, y_train):
    X_train_k, X_val_k = X_train[train_idx], X_train[val_idx]
    y_train_k, y_val_k = y_train.iloc[train_idx],
y_train.iloc[val_idx]

    nb_model = GaussianNB()
    nb_model.fit(X_train_k, y_train_k)
    y_pred_k = nb_model.predict(X_val_k)
    y_prob_k = nb_model.predict_proba(X_val_k)[: , 1]

    metrics = calculate_metrics(y_val_k, y_pred_k, y_prob_k)
    nb_metrics.append(metrics)

nb_metrics_df = pd.DataFrame(nb_metrics)
print("\nNaive Bayes Metrics Per Fold:")
display(nb_metrics_df)

# Average Metrics for Naive Bayes
nb_avg_metrics = nb_metrics_df.mean()
print("\nNaive Bayes Average Metrics:")
print(nb_avg_metrics)
```

Naive Bayes Metrics Per Fold:

	TP	TN	FP	FN	P	N	TPR	TNR	FPR
FNR \									
0	1700	1062	887	674	2374	1949	0.716091	0.544895	0.455105
0.283909									
1	1705	1030	919	669	2374	1949	0.718197	0.528476	0.471524
0.281803									
2	1697	1049	900	677	2374	1949	0.714827	0.538225	0.461775
0.285173									
3	1700	1066	883	674	2374	1949	0.716091	0.546947	0.453053
0.283909									
4	1658	1051	898	716	2374	1949	0.698399	0.539251	0.460749
0.301601									
5	1697	1044	905	677	2374	1949	0.714827	0.535659	0.464341
0.285173									
6	1704	1053	896	670	2374	1949	0.717776	0.540277	0.459723
0.282224									
7	1713	1064	885	661	2374	1949	0.721567	0.545921	0.454079
0.278433									
8	1680	1055	893	694	2374	1948	0.707666	0.541581	0.458419
0.292334									
9	1685	1018	930	689	2374	1948	0.709773	0.522587	0.477413
0.290227									
	Precision		F1	Accuracy	Error Rate	Balanced Accuracy			
TSS \									
0	0.657132	0.685346	0.638908	0.361092		0.630493			
0.260986									
1	0.649771	0.682273	0.632663	0.367337		0.623337			
0.246673									
2	0.653446	0.682760	0.635207	0.364793		0.626526			
0.253052									
3	0.658149	0.685899	0.639833	0.360167		0.631519			
0.263038									
4	0.648670	0.672617	0.626648	0.373352		0.618825			
0.237650									
5	0.652191	0.682074	0.634050	0.365950		0.625243			
0.250487									
6	0.655385	0.685163	0.637752	0.362248		0.629026			
0.258053									
7	0.659353	0.689059	0.642378	0.357622		0.633744			
0.267488									
8	0.652934	0.679200	0.632809	0.367191		0.624624			
0.249247									
9	0.644359	0.675486	0.625405	0.374595		0.616180			
0.232360									
	HSS	Brier Score	Brier Skill Score	ROC AUC					
0	0.260986	0.225493	0.089226	0.698857					
1	0.246673	0.231272	0.065882	0.684202					
2	0.253052	0.230491	0.069038	0.689188					

3	0.263038	0.228534	0.076942	0.692784
4	0.237650	0.234949	0.051031	0.679107
5	0.250487	0.234039	0.054708	0.683308
6	0.258053	0.228622	0.076587	0.694074
7	0.267488	0.227510	0.081077	0.695671
8	0.249247	0.231696	0.064123	0.688908
9	0.232360	0.237739	0.039715	0.673114

Naive Bayes Average Metrics:

TP	1693.900000
TN	1049.200000
FP	899.600000
FN	680.100000
P	2374.000000
N	1948.800000
TPR	0.713521
TNR	0.538382
FPR	0.461618
FNR	0.286479
Precision	0.653139
F1	0.681988
Accuracy	0.634565
Error Rate	0.365435
Balanced Accuracy	0.625952
TSS	0.251903
HSS	0.251903
Brier Score	0.231035
Brier Skill Score	0.066833
ROC AUC	0.687921
dtype: float64	

## K-Nearest Neighbors (KNN) Classification with K-Fold Cross-Validation

In this cell, we train and evaluate a **K-Nearest Neighbors (KNN)** model using **10-Fold Cross-Validation**. This approach splits the dataset into 10 subsets (folds) and evaluates the model's performance on each fold to ensure robust results.

### Process

- Data Preparation:**
  - Convert `X_train` and `y_train` into NumPy arrays to facilitate efficient slicing and manipulation for K-Fold splits.
- K-Fold Cross-Validation:**
  - The dataset is split into 10 folds using `KFold`, with shuffling enabled to randomize data distribution across folds.

- For each fold:
  - Training and validation subsets are created.
  - A **K-Nearest Neighbors (KNN)** classifier is initialized with `n_neighbors=5`.
  - The model is trained on the training subset.
  - Predictions (`y_pred_k`) and probability scores (`y_prob_k`) are generated for the validation subset.
- 3. **Metric Calculation:**
  - Metrics are computed for each fold using the custom `calculate_metrics` function, which evaluates:
    - **Accuracy**
    - **Precision**
    - **Recall**
    - **F1-Score**
    - **Brier Score**
    - And other derived metrics
  - Metrics for each fold are stored in a list.
- 4. **Metrics Aggregation:**
  - The metrics for all folds are stored in a DataFrame for detailed analysis.
  - The average values of the metrics are computed to summarize the overall performance of the KNN classifier.

## Output

- **Fold-wise Metrics:** Displays performance metrics for each fold, providing a fold-by-fold breakdown of the KNN model's performance.
- **Average Metrics:** Presents the mean values of all metrics across the 10 folds, offering an overall view of the KNN classifier's strengths and weaknesses.

## Observations

The **KNN model** uses the proximity of data points to make predictions, which can work well with well-separated classes but may struggle with high-dimensional or imbalanced data. The evaluation includes standard classification metrics as well as probability calibration metrics like the **Brier Score**, providing a holistic assessment of the model's performance.

```
# KNN Model
knn_metrics = []
for train_idx, val_idx in kf.split(X_train, y_train):
    X_train_k, X_val_k = X_train[train_idx], X_train[val_idx]
    y_train_k, y_val_k = y_train.iloc[train_idx],
y_train.iloc[val_idx]

    knn_model = KNeighborsClassifier(n_neighbors=5)
    knn_model.fit(X_train_k, y_train_k)
    y_pred_k = knn_model.predict(X_val_k)
    y_prob_k = knn_model.predict_proba(X_val_k)[: , 1]
```

```

metrics = calculate_metrics(y_val_k, y_pred_k, y_prob_k)
knn_metrics.append(metrics)

knn_metrics_df = pd.DataFrame(knn_metrics)
print("\nKNN Metrics Per Fold:")
display(knn_metrics_df)

# Average Metrics for KNN
knn_avg_metrics = knn_metrics_df.mean()
print("\nKNN Average Metrics:")
print(knn_avg_metrics)

```

KNN Metrics Per Fold:

	TP	TN	FP	FN	P	N	TPR	TNR	FPR
FNR \									
0	2354	1915	34	20	2374	1949	0.991575	0.982555	0.017445
0.008425									
1	2350	1926	23	24	2374	1949	0.989890	0.988199	0.011801
0.010110									
2	2358	1920	29	16	2374	1949	0.993260	0.985121	0.014879
0.006740									
3	2354	1917	32	20	2374	1949	0.991575	0.983581	0.016419
0.008425									
4	2343	1916	33	31	2374	1949	0.986942	0.983068	0.016932
0.013058									
5	2348	1912	37	26	2374	1949	0.989048	0.981016	0.018984
0.010952									
6	2350	1915	34	24	2374	1949	0.989890	0.982555	0.017445
0.010110									
7	2353	1910	39	21	2374	1949	0.991154	0.979990	0.020010
0.008846									
8	2355	1922	26	19	2374	1948	0.991997	0.986653	0.013347
0.008003									
9	2358	1918	30	16	2374	1948	0.993260	0.984600	0.015400
0.006740									

	Precision	F1	Accuracy	Error Rate	Balanced Accuracy
TSS \					
0	0.985762	0.988660	0.987509	0.012491	0.987065
0.974131					
1	0.990308	0.990099	0.989128	0.010872	0.989045
0.978090					
2	0.987851	0.990548	0.989591	0.010409	0.989190
0.978381					
3	0.986588	0.989076	0.987971	0.012029	0.987578
0.975157					
4	0.986111	0.986526	0.985195	0.014805	0.985005
0.970010					

5	0.984486	0.986762	0.985427	0.014573	0.985032
0.970064					
6	0.985738	0.987810	0.986583	0.013417	0.986223
0.972446					
7	0.983696	0.987411	0.986121	0.013879	0.985572
0.971144					
8	0.989080	0.990536	0.989588	0.010412	0.989325
0.978650					
9	0.987437	0.990340	0.989357	0.010643	0.988930
0.977860					

	HSS	Brier Score	Brier Skill Score	ROC AUC
0	0.974131	0.009308	0.962403	0.997212
1	0.978090	0.007689	0.968943	0.998696
2	0.978381	0.008614	0.965206	0.997452
3	0.975157	0.009114	0.963188	0.997698
4	0.970010	0.010493	0.957620	0.997352
5	0.970064	0.011298	0.954368	0.996292
6	0.972446	0.010557	0.957358	0.996080
7	0.971144	0.011159	0.954929	0.995872
8	0.978650	0.007987	0.967738	0.996801
9	0.977860	0.008700	0.964860	0.997733

#### KNN Average Metrics:

TP	2352.300000
TN	1917.100000
FP	31.700000
FN	21.700000
P	2374.000000
N	1948.800000
TPR	0.990859
TNR	0.983734
FPR	0.016266
FNR	0.009141
Precision	0.986706
F1	0.988777
Accuracy	0.987647
Error Rate	0.012353
Balanced Accuracy	0.987297
TSS	0.974593
HSS	0.974593
Brier Score	0.009492
Brier Skill Score	0.961661
ROC AUC	0.997119
dtype:	float64



# Long Short-Term Memory (LSTM) Classification with K-Fold Cross-Validation

We train and evaluate a **Long Short-Term Memory (LSTM)** deep learning model using **10-Fold Cross-Validation**. This method ensures robust evaluation by splitting the dataset into 10 subsets (folds) and rotating the training and validation sets across folds.

## Process

### 1. Data Preparation:

- The input features are reshaped to match the LSTM model's expected input format (3D: samples, timesteps, features).
- `X_train` and `X_test` are expanded along the last dimension to prepare the data for sequential processing by LSTM layers.

### 2. LSTM Model Design:

- The LSTM architecture includes:
  - A **Bidirectional LSTM** layer to capture both forward and backward dependencies in the data.
  - **Dropout layers** to reduce overfitting.
  - A **Dense layer** with a sigmoid activation function for binary classification.
- The model is compiled using the **Adam optimizer** and **binary cross-entropy loss**.

### 3. K-Fold Cross-Validation:

- The dataset is split into 10 folds using `StratifiedKFold`, ensuring class balance in each fold.
- For each fold:
  - Training and validation subsets are created.
  - The LSTM model is trained using **early stopping** to prevent overfitting.
  - Predictions (`y_pred_k`) and probability scores (`y_prob_k`) are generated for the validation subset.

### 4. Metric Calculation:

- Performance metrics are computed for each fold using the `calculate_metrics` function. These include:
  - **Accuracy**
  - **Precision**
  - **Recall**
  - **F1-Score**
  - **Brier Score**
  - And other derived metrics
- Metrics for each fold are stored in a list for analysis.

### 5. Metrics Aggregation:

- The metrics for all folds are consolidated into a DataFrame.
- The average values of the metrics are calculated to summarize the overall performance of the LSTM model.

## Output

- **Fold-wise Metrics:** Displays detailed performance metrics for each fold, showing the variability in the LSTM model's performance across different subsets of the data.
- **Average Metrics:** Presents the mean values of all metrics across the 10 folds, providing a comprehensive evaluation of the LSTM model.

## Observations

The **LSTM model** is capable of capturing complex sequential patterns and relationships in the data. It performs particularly well when there is temporal or positional significance in the features. The evaluation includes standard classification metrics and probability calibration metrics like the **Brier Score**, ensuring a well-rounded assessment of the model's effectiveness.

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense,
Bidirectional, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
import numpy as np
from sklearn.metrics import roc_curve, auc

# Prepare data for LSTM (reshape X_train and X_test)
X_train_lstm = np.expand_dims(X_train, axis=-1)
X_test_lstm = np.expand_dims(X_test, axis=-1)

# Define LSTM model
def create_lstm_model(input_shape):
    model = Sequential()
    model.add(Bidirectional(LSTM(64, return_sequences=True,
input_shape=input_shape)))
    model.add(Dropout(0.5))
    model.add(LSTM(32, return_sequences=False))
    model.add(Dropout(0.5))
    model.add(Dense(1, activation='sigmoid'))
    model.compile(optimizer=Adam(learning_rate=0.001),
loss='binary_crossentropy', metrics=['accuracy'])
    return model

# Cross-validation for LSTM
kf = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)
lstm_metrics = []

print("Starting LSTM Cross-Validation...\n")
for fold, (train_idx, val_idx) in enumerate(kf.split(X_train,
y_train), start=1):
    print(f"Training Fold {fold}...")

    # Split the data for this fold
    X_train_k, X_val_k = X_train_lstm[train_idx],
```

```

X_train_lstm[val_idx]
    y_train_k, y_val_k = y_train.iloc[train_idx],
y_train.iloc[val_idx]

    # Create and compile LSTM model
    lstm_model = create_lstm_model((X_train_k.shape[1], 1))

    # Early stopping
    early_stopping = EarlyStopping(monitor='val_loss', patience=3,
restore_best_weights=True)

    # Train the model
    lstm_model.fit(X_train_k, y_train_k, validation_data=(X_val_k,
y_val_k),
                    epochs=5, batch_size=64,
callbacks=[early_stopping], verbose=0)

    # Predictions
    y_prob_k = lstm_model.predict(X_val_k).flatten()
    y_pred_k = (y_prob_k > 0.5).astype(int)

    # Calculate metrics
    metrics = calculate_metrics(y_val_k, y_pred_k, y_prob_k)
    lstm_metrics.append(metrics)

    # Print metrics for this fold
    print(f"Fold {fold} Metrics:")
    for metric_name, value in metrics.items():
        print(f"    - {metric_name}: {value:.4f}")
    print()

# Convert metrics to DataFrame for analysis
lstm_metrics_df = pd.DataFrame(lstm_metrics)
print("\nLSTM Metrics Per Fold:")
display(lstm_metrics_df)

# Average Metrics for LSTM
lstm_avg_metrics = lstm_metrics_df.mean()
print("\nLSTM Average Metrics:")
for metric_name, value in lstm_avg_metrics.items():
    print(f"    - {metric_name}: {value:.4f}")

# Compare All Models
print("\n----- Iteration-Wise Metrics Comparison -----")
algorithms = {
    "Random Forest": rf_metrics_df,
    "Naive Bayes": nb_metrics_df,
    "KNN": knn_metrics_df,
    "LSTM": lstm_metrics_df
}

```

```

for iteration in range(len(knn_metrics_df)):
    print(f"Iteration {iteration + 1}:")
    print("----- Metrics for all Algorithms in Iteration -----")
    comparison_table = {}
    for algo_name, metrics_df in algorithms.items():
        comparison_table[algo_name] = metrics_df.iloc[iteration]

    # Convert to DataFrame for display
    iteration_comparison_df = pd.DataFrame(comparison_table).T
    iteration_comparison_df.columns = rf_metrics_df.columns # Use
metric names as columns
    display(iteration_comparison_df)

# Compute average metrics for each algorithm
avg_metrics = {}
for algo_name, metrics_df in algorithms.items():
    avg_metrics[algo_name] = metrics_df.mean()

# Convert to DataFrame for display
avg_metrics_df = pd.DataFrame(avg_metrics).T
avg_metrics_df.columns = rf_metrics_df.columns # Use metric names as
columns
print("\n----- Average Metrics Across All Iterations ----- \n")
display(avg_metrics_df)

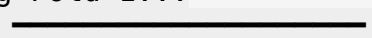
# Combine average metrics for all algorithms
results_df = avg_metrics_df[['Accuracy', 'F1', 'Precision', 'ROC
AUC']]
print("\nComparison of Average Metrics Across Models:")
display(results_df)

# Plot results
results_df.plot(kind='bar', figsize=(12, 8))
plt.title("Comparison of Average Metrics Across Models")
plt.ylabel("Scores")
plt.xticks(rotation=45)
plt.legend(loc="upper right")
plt.show()

```

Starting LSTM Cross-Validation...

Training Fold 1...

136/136  0s 3ms/step

Fold 1 Metrics:

- TP: 1954.0000
- TN: 1549.0000
- FP: 400.0000
- FN: 420.0000
- P: 2374.0000

- N: 1949.0000
- TPR: 0.8231
- TNR: 0.7948
- FPR: 0.2052
- FNR: 0.1769
- Precision: 0.8301
- F1: 0.8266
- Accuracy: 0.8103
- Error Rate: 0.1897
- Balanced Accuracy: 0.8089
- TSS: 0.6178
- HSS: 0.6178
- Brier Score: 0.1297
- Brier Skill Score: 0.4760
- ROC AUC: 0.8988

Training Fold 2...

136/136 ————— 0s 2ms/step

Fold 2 Metrics:

- TP: 1843.0000
- TN: 1616.0000
- FP: 333.0000
- FN: 531.0000
- P: 2374.0000
- N: 1949.0000
- TPR: 0.7763
- TNR: 0.8291
- FPR: 0.1709
- FNR: 0.2237
- Precision: 0.8470
- F1: 0.8101
- Accuracy: 0.8001
- Error Rate: 0.1999
- Balanced Accuracy: 0.8027
- TSS: 0.6055
- HSS: 0.6055
- Brier Score: 0.1417
- Brier Skill Score: 0.4277
- ROC AUC: 0.8808

Training Fold 3...

136/136 ————— 1s 2ms/step

Fold 3 Metrics:

- TP: 1890.0000
- TN: 1465.0000
- FP: 484.0000
- FN: 484.0000
- P: 2374.0000
- N: 1949.0000

- TPR: 0.7961
- TNR: 0.7517
- FPR: 0.2483
- FNR: 0.2039
- Precision: 0.7961
- F1: 0.7961
- Accuracy: 0.7761
- Error Rate: 0.2239
- Balanced Accuracy: 0.7739
- TSS: 0.5478
- HSS: 0.5478
- Brier Score: 0.1519
- Brier Skill Score: 0.3865
- ROC AUC: 0.8601

Training Fold 4...

136/136 ————— 0s 2ms/step

Fold 4 Metrics:

- TP: 1822.0000
- TN: 1549.0000
- FP: 400.0000
- FN: 552.0000
- P: 2374.0000
- N: 1949.0000
- TPR: 0.7675
- TNR: 0.7948
- FPR: 0.2052
- FNR: 0.2325
- Precision: 0.8200
- F1: 0.7929
- Accuracy: 0.7798
- Error Rate: 0.2202
- Balanced Accuracy: 0.7811
- TSS: 0.5622
- HSS: 0.5622
- Brier Score: 0.1502
- Brier Skill Score: 0.3933
- ROC AUC: 0.8624

Training Fold 5...

136/136 ————— 0s 2ms/step

Fold 5 Metrics:

- TP: 1965.0000
- TN: 1545.0000
- FP: 404.0000
- FN: 409.0000
- P: 2374.0000
- N: 1949.0000
- TPR: 0.8277

- TNR: 0.7927
- FPR: 0.2073
- FNR: 0.1723
- Precision: 0.8295
- F1: 0.8286
- Accuracy: 0.8119
- Error Rate: 0.1881
- Balanced Accuracy: 0.8102
- TSS: 0.6204
- HSS: 0.6204
- Brier Score: 0.1312
- Brier Skill Score: 0.4700
- ROC AUC: 0.8954

Training Fold 6...

136/136 ————— 0s 2ms/step

Fold 6 Metrics:

- TP: 1922.0000
- TN: 1530.0000
- FP: 419.0000
- FN: 452.0000
- P: 2374.0000
- N: 1949.0000
- TPR: 0.8096
- TNR: 0.7850
- FPR: 0.2150
- FNR: 0.1904
- Precision: 0.8210
- F1: 0.8153
- Accuracy: 0.7985
- Error Rate: 0.2015
- Balanced Accuracy: 0.7973
- TSS: 0.5946
- HSS: 0.5946
- Brier Score: 0.1394
- Brier Skill Score: 0.4368
- ROC AUC: 0.8852

Training Fold 7...

136/136 ————— 0s 2ms/step

Fold 7 Metrics:

- TP: 1944.0000
- TN: 1489.0000
- FP: 460.0000
- FN: 430.0000
- P: 2374.0000
- N: 1949.0000
- TPR: 0.8189
- TNR: 0.7640

- FPR: 0.2360
- FNR: 0.1811
- Precision: 0.8087
- F1: 0.8137
- Accuracy: 0.7941
- Error Rate: 0.2059
- Balanced Accuracy: 0.7914
- TSS: 0.5829
- HSS: 0.5829
- Brier Score: 0.1468
- Brier Skill Score: 0.4069
- ROC AUC: 0.8724

Training Fold 8...

136/136 ————— 0s 2ms/step

Fold 8 Metrics:

- TP: 2038.0000
- TN: 1414.0000
- FP: 535.0000
- FN: 336.0000
- P: 2374.0000
- N: 1949.0000
- TPR: 0.8585
- TNR: 0.7255
- FPR: 0.2745
- FNR: 0.1415
- Precision: 0.7921
- F1: 0.8239
- Accuracy: 0.7985
- Error Rate: 0.2015
- Balanced Accuracy: 0.7920
- TSS: 0.5840
- HSS: 0.5840
- Brier Score: 0.1417
- Brier Skill Score: 0.4277
- ROC AUC: 0.8796

Training Fold 9...

136/136 ————— 0s 2ms/step

Fold 9 Metrics:

- TP: 2009.0000
- TN: 1440.0000
- FP: 508.0000
- FN: 365.0000
- P: 2374.0000
- N: 1948.0000
- TPR: 0.8463
- TNR: 0.7392
- FPR: 0.2608



- FNR: 0.1537
- Precision: 0.7982
- F1: 0.8215
- Accuracy: 0.7980
- Error Rate: 0.2020
- Balanced Accuracy: 0.7927
- TSS: 0.5855
- HSS: 0.5855
- Brier Score: 0.1367
- Brier Skill Score: 0.4479
- ROC AUC: 0.8910

Training Fold 10...

136/136 ————— 0s 2ms/step

Fold 10 Metrics:

- TP: 1814.0000
- TN: 1558.0000
- FP: 390.0000
- FN: 560.0000
- P: 2374.0000
- N: 1948.0000
- TPR: 0.7641
- TNR: 0.7998
- FPR: 0.2002
- FNR: 0.2359
- Precision: 0.8230
- F1: 0.7925
- Accuracy: 0.7802
- Error Rate: 0.2198
- Balanced Accuracy: 0.7820
- TSS: 0.5639
- HSS: 0.5639
- Brier Score: 0.1416
- Brier Skill Score: 0.4279
- ROC AUC: 0.8787

LSTM Metrics Per Fold:

	TP	TN	FP	FN	P	N	TPR	TNR	FPR
FNR \									
0	1954	1549	400	420	2374	1949	0.823083	0.794767	0.205233
0.176917									
1	1843	1616	333	531	2374	1949	0.776327	0.829143	0.170857
0.223673									
2	1890	1465	484	484	2374	1949	0.796125	0.751668	0.248332
0.203875									
3	1822	1549	400	552	2374	1949	0.767481	0.794767	0.205233
0.232519									
4	1965	1545	404	409	2374	1949	0.827717	0.792714	0.207286

0.172283  
5 1922 1530 419 452 2374 1949 0.809604 0.785018 0.214982  
0.190396  
6 1944 1489 460 430 2374 1949 0.818871 0.763982 0.236018  
0.181129  
7 2038 1414 535 336 2374 1949 0.858467 0.725500 0.274500  
0.141533  
8 2009 1440 508 365 2374 1948 0.846251 0.739220 0.260780  
0.153749  
9 1814 1558 390 560 2374 1948 0.764111 0.799795 0.200205  
0.235889

	Precision	F1	Accuracy	Error Rate	Balanced Accuracy
TSS \					
0	0.830076	0.826565	0.810317	0.189683	0.808925
0.617850					
1	0.846967	0.810110	0.800139	0.199861	0.802735
0.605470					
2	0.796125	0.796125	0.776081	0.223919	0.773896
0.547792					
3	0.819982	0.792863	0.779783	0.220217	0.781124
0.562248					
4	0.829464	0.828590	0.811936	0.188064	0.810216
0.620431					
5	0.821017	0.815270	0.798520	0.201480	0.797311
0.594622					
6	0.808652	0.813730	0.794124	0.205876	0.791426
0.582853					
7	0.792072	0.823934	0.798520	0.201480	0.791983
0.583967					
8	0.798172	0.821509	0.798010	0.201990	0.792735
0.585471					
9	0.823049	0.792486	0.780194	0.219806	0.781953
0.563906					

	HSS	Brier Score	Brier Skill Score	ROC AUC
0	0.617850	0.129737	0.475989	0.898849
1	0.605470	0.141682	0.427739	0.880773
2	0.547792	0.151899	0.386473	0.860094
3	0.562248	0.150197	0.393349	0.862358
4	0.620431	0.131229	0.469960	0.895385
5	0.594622	0.139439	0.436801	0.885173
6	0.582853	0.146844	0.406892	0.872436
7	0.583967	0.141703	0.427658	0.879604
8	0.585471	0.136685	0.447897	0.890977
9	0.563906	0.141638	0.427888	0.878685

LSTM Average Metrics:  
- TP: 1920.1000

- TN: 1515.5000
- FP: 433.3000
- FN: 453.9000
- P: 2374.0000
- N: 1948.8000
- TPR: 0.8088
- TNR: 0.7777
- FPR: 0.2223
- FNR: 0.1912
- Precision: 0.8166
- F1: 0.8121
- Accuracy: 0.7948
- Error Rate: 0.2052
- Balanced Accuracy: 0.7932
- TSS: 0.5865
- HSS: 0.5865
- Brier Score: 0.1411
- Brier Skill Score: 0.4301
- ROC AUC: 0.8804

#### ----- Iteration-Wise Metrics Comparison -----

Iteration 1:

#### ----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2351.0	1925.0	24.0	23.0	2374.0	1949.0	0.990312
Naive Bayes	1700.0	1062.0	887.0	674.0	2374.0	1949.0	0.716091
KNN	2354.0	1915.0	34.0	20.0	2374.0	1949.0	0.991575
LSTM	1954.0	1549.0	400.0	420.0	2374.0	1949.0	0.823083

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest	0.987686	0.012314	0.009688	0.989895	0.990103
0.989128					
Naive Bayes	0.544895	0.455105	0.283909	0.657132	0.685346
0.638908					
KNN	0.982555	0.017445	0.008425	0.985762	0.988660
0.987509					
LSTM	0.794767	0.205233	0.176917	0.830076	0.826565
0.810317					

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest	0.010872	0.988999	0.977998	0.977998

0.010200				
Naive Bayes	0.361092	0.630493	0.260986	0.260986
0.225493				
KNN	0.012491	0.987065	0.974131	0.974131
0.009308				
LSTM	0.189683	0.808925	0.617850	0.617850
0.129737				

	Brier Skill Score	ROC AUC
Random Forest	0.958803	0.999480
Naive Bayes	0.089226	0.698857
KNN	0.962403	0.997212
LSTM	0.475989	0.898849

Iteration 2:

----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2352.0	1935.0	14.0	22.0	2374.0	1949.0	0.990733
Naive Bayes	1705.0	1030.0	919.0	669.0	2374.0	1949.0	0.718197
KNN	2350.0	1926.0	23.0	24.0	2374.0	1949.0	0.989890
LSTM	1843.0	1616.0	333.0	531.0	2374.0	1949.0	0.776327

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest	0.992817	0.007183	0.009267	0.994083	0.992405
0.991672					
Naive Bayes	0.528476	0.471524	0.281803	0.649771	0.682273
0.632663					
KNN	0.988199	0.011801	0.010110	0.990308	0.990099
0.989128					
LSTM	0.829143	0.170857	0.223673	0.846967	0.810110
0.800139					

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest	0.008328	0.991775	0.983550	0.983550
0.008712				
Naive Bayes	0.367337	0.623337	0.246673	0.246673
0.231272				
KNN	0.010872	0.989045	0.978090	0.978090
0.007689				
LSTM	0.199861	0.802735	0.605470	0.605470
0.141682				

	Brier Skill Score	ROC AUC
Random Forest	0.964814	0.999620
Naive Bayes	0.065882	0.684202
KNN	0.968943	0.998696
LSTM	0.427739	0.880773

Iteration 3:

----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2353.0	1924.0	25.0	21.0	2374.0	1949.0	0.991154
Naive Bayes	1697.0	1049.0	900.0	677.0	2374.0	1949.0	0.714827
KNN	2358.0	1920.0	29.0	16.0	2374.0	1949.0	0.993260
LSTM	1890.0	1465.0	484.0	484.0	2374.0	1949.0	0.796125

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest	0.987173	0.012827	0.008846	0.989487	0.990320
0.989359					
Naive Bayes	0.538225	0.461775	0.285173	0.653446	0.682760
0.635207					
KNN	0.985121	0.014879	0.006740	0.987851	0.990548
0.989591					
LSTM	0.751668	0.248332	0.203875	0.796125	0.796125
0.776081					

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest	0.010641	0.989164	0.978327	0.978327
0.009757				
Naive Bayes	0.364793	0.626526	0.253052	0.253052
0.230491				
KNN	0.010409	0.989190	0.978381	0.978381
0.008614				
LSTM	0.223919	0.773896	0.547792	0.547792
0.151899				

	Brier Skill Score	ROC AUC
Random Forest	0.960590	0.999516
Naive Bayes	0.069038	0.689188
KNN	0.965206	0.997452
LSTM	0.386473	0.860094

Iteration 4:

----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\ Random Forest	2354.0	1925.0	24.0	20.0	2374.0	1949.0	0.991575
Naive Bayes	1700.0	1066.0	883.0	674.0	2374.0	1949.0	0.716091
KNN	2354.0	1917.0	32.0	20.0	2374.0	1949.0	0.991575
LSTM	1822.0	1549.0	400.0	552.0	2374.0	1949.0	0.767481
	TNR	FPR	FNR	Precision	F1		
Accuracy \							
Random Forest	0.987686	0.012314	0.008425	0.989907	0.990741		
0.989822							
Naive Bayes	0.546947	0.453053	0.283909	0.658149	0.685899		
0.639833							
KNN	0.983581	0.016419	0.008425	0.986588	0.989076		
0.987971							
LSTM	0.794767	0.205233	0.232519	0.819982	0.792863		
0.779783							
	Error Rate	Balanced Accuracy	TSS	HSS			
Brier Score \							
Random Forest	0.010178		0.989631	0.979261	0.979261		
0.010139							
Naive Bayes	0.360167		0.631519	0.263038	0.263038		
0.228534							
KNN	0.012029		0.987578	0.975157	0.975157		
0.009114							
LSTM	0.220217		0.781124	0.562248	0.562248		
0.150197							
	Brier Skill Score	ROC AUC					
Random Forest	0.959048	0.999159					
Naive Bayes	0.076942	0.692784					
KNN	0.963188	0.997698					
LSTM	0.393349	0.862358					
Iteration 5:							
----- Metrics for all Algorithms in Iteration -----							
	TP	TN	FP	FN	P	N	TPR
\ Random Forest	2345.0	1924.0	25.0	29.0	2374.0	1949.0	0.987784
Naive Bayes	1658.0	1051.0	898.0	716.0	2374.0	1949.0	0.698399
KNN	2343.0	1916.0	33.0	31.0	2374.0	1949.0	0.986942
LSTM	1965.0	1545.0	404.0	409.0	2374.0	1949.0	0.827717

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest	0.987173	0.012827	0.012216	0.989451	0.988617
0.987509					
Naive Bayes	0.539251	0.460749	0.301601	0.648670	0.672617
0.626648					
KNN	0.983068	0.016932	0.013058	0.986111	0.986526
0.985195					
LSTM	0.792714	0.207286	0.172283	0.829464	0.828590
0.811936					

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest	0.012491	0.987479	0.974957	0.974957
0.010940				
Naive Bayes	0.373352	0.618825	0.237650	0.237650
0.234949				
KNN	0.014805	0.985005	0.970010	0.970010
0.010493				
LSTM	0.188064	0.810216	0.620431	0.620431
0.131229				

	Brier Skill Score	ROC AUC
Random Forest	0.955813	0.998922
Naive Bayes	0.051031	0.679107
KNN	0.957620	0.997352
LSTM	0.469960	0.895385

Iteration 6:  
 ----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2343.0	1931.0	18.0	31.0	2374.0	1949.0	0.986942
Naive Bayes	1697.0	1044.0	905.0	677.0	2374.0	1949.0	0.714827
KNN	2348.0	1912.0	37.0	26.0	2374.0	1949.0	0.989048
LSTM	1922.0	1530.0	419.0	452.0	2374.0	1949.0	0.809604

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest	0.990764	0.009236	0.013058	0.992376	0.989652
0.988665					
Naive Bayes	0.535659	0.464341	0.285173	0.652191	0.682074
0.634050					
KNN	0.981016	0.018984	0.010952	0.984486	0.986762

0.985427						
LSTM	0.785018	0.214982	0.190396	0.821017	0.815270	
0.798520						

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest	0.011335	0.988853	0.977706	0.977706
0.011084				
Naive Bayes	0.365950	0.625243	0.250487	0.250487
0.234039				
KNN	0.014573	0.985032	0.970064	0.970064
0.011298				
LSTM	0.201480	0.797311	0.594622	0.594622
0.139439				

	Brier Skill Score	ROC AUC
Random Forest	0.955233	0.998800
Naive Bayes	0.054708	0.683308
KNN	0.954368	0.996292
LSTM	0.436801	0.885173

Iteration 7:

----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2348.0	1927.0	22.0	26.0	2374.0	1949.0	0.989048
Naive Bayes	1704.0	1053.0	896.0	670.0	2374.0	1949.0	0.717776
KNN	2350.0	1915.0	34.0	24.0	2374.0	1949.0	0.989890
LSTM	1944.0	1489.0	460.0	430.0	2374.0	1949.0	0.818871

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest	0.988712	0.011288	0.010952	0.990717	0.989882
0.988897					
Naive Bayes	0.540277	0.459723	0.282224	0.655385	0.685163
0.637752					
KNN	0.982555	0.017445	0.010110	0.985738	0.987810
0.986583					
LSTM	0.763982	0.236018	0.181129	0.808652	0.813730
0.794124					

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest	0.011103	0.988880	0.977760	0.977760
0.010409				



Naive Bayes 0.228622	0.362248	0.629026	0.258053	0.258053
KNN 0.010557	0.013417	0.986223	0.972446	0.972446
LSTM 0.146844	0.205876	0.791426	0.582853	0.582853

	Brier Skill Score	ROC AUC
Random Forest	0.957960	0.999214
Naive Bayes	0.076587	0.694074
KNN	0.957358	0.996080
LSTM	0.406892	0.872436

Iteration 8:

----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2350.0	1930.0	19.0	24.0	2374.0	1949.0	0.989890
Naive Bayes	1713.0	1064.0	885.0	661.0	2374.0	1949.0	0.721567
KNN	2353.0	1910.0	39.0	21.0	2374.0	1949.0	0.991154
LSTM	2038.0	1414.0	535.0	336.0	2374.0	1949.0	0.858467

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest 0.990053	0.990251	0.009749	0.010110	0.991980	0.990934
Naive Bayes 0.642378	0.545921	0.454079	0.278433	0.659353	0.689059
KNN 0.986121	0.979990	0.020010	0.008846	0.983696	0.987411
LSTM 0.798520	0.725500	0.274500	0.141533	0.792072	0.823934

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest 0.010718	0.009947	0.990071	0.980142	0.980142
Naive Bayes 0.227510	0.357622	0.633744	0.267488	0.267488
KNN 0.011159	0.013879	0.985572	0.971144	0.971144
LSTM 0.141703	0.201480	0.791983	0.583967	0.583967

	Brier Skill Score	ROC AUC
--	-------------------	---------

Random Forest	0.956709	0.999359
Naive Bayes	0.081077	0.695671
KNN	0.954929	0.995872
LSTM	0.427658	0.879604

Iteration 9:

----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\ Random Forest	2357.0	1931.0	17.0	17.0	2374.0	1948.0	0.992839
Naive Bayes	1680.0	1055.0	893.0	694.0	2374.0	1948.0	0.707666
KNN	2355.0	1922.0	26.0	19.0	2374.0	1948.0	0.991997
LSTM	2009.0	1440.0	508.0	365.0	2374.0	1948.0	0.846251

	TNR	FPR	FNR	Precision	F1
Accuracy \					
Random Forest	0.991273	0.008727	0.007161	0.992839	0.992839
0.992133					
Naive Bayes	0.541581	0.458419	0.292334	0.652934	0.679200
0.632809					
KNN	0.986653	0.013347	0.008003	0.989080	0.990536
0.989588					
LSTM	0.739220	0.260780	0.153749	0.798172	0.821509
0.798010					

	Error Rate	Balanced Accuracy	TSS	HSS
Brier Score \				
Random Forest	0.007867	0.992056	0.984112	0.984112
0.009561				
Naive Bayes	0.367191	0.624624	0.249247	0.249247
0.231696				
KNN	0.010412	0.989325	0.978650	0.978650
0.007987				
LSTM	0.201990	0.792735	0.585471	0.585471
0.136685				


	Brier Skill Score	ROC AUC
Random Forest	0.961382	0.999157
Naive Bayes	0.064123	0.688908
KNN	0.967738	0.996801
LSTM	0.447897	0.890977

Iteration 10:

----- Metrics for all Algorithms in Iteration -----

	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2355.0	1926.0	22.0	19.0	2374.0	1948.0	0.991997
Naive Bayes	1685.0	1018.0	930.0	689.0	2374.0	1948.0	0.709773
KNN	2358.0	1918.0	30.0	16.0	2374.0	1948.0	0.993260
LSTM	1814.0	1558.0	390.0	560.0	2374.0	1948.0	0.764111
	TNR	FPR	FNR	Precision	F1		
Accuracy \							
Random Forest	0.988706	0.011294	0.008003	0.990745	0.991370	0.990514	
Naive Bayes	0.522587	0.477413	0.290227	0.644359	0.675486	0.625405	
KNN	0.984600	0.015400	0.006740	0.987437	0.990340	0.989357	
LSTM	0.799795	0.200205	0.235889	0.823049	0.792486	0.780194	
	Error Rate	Balanced Accuracy	TSS		HSS		
Brier Score \							
Random Forest	0.009486	0.990351	0.980703	0.980703	0.009352	0.980703	
Naive Bayes	0.374595	0.616180	0.232360	0.232360	0.237739	0.232360	
KNN	0.010643	0.988930	0.977860	0.977860	0.008700	0.977860	
LSTM	0.219806	0.781953	0.563906	0.563906	0.141638	0.563906	
	Brier Skill Score	ROC AUC					
Random Forest	0.962224	0.999436					
Naive Bayes	0.039715	0.673114					
KNN	0.964860	0.997733					
LSTM	0.427888	0.878685					
----- Average Metrics Across All Iterations -----							
	TP	TN	FP	FN	P	N	TPR
\							
Random Forest	2350.8	1927.8	21.0	23.2	2374.0	1948.8	0.990227
Naive Bayes	1693.9	1049.2	899.6	680.1	2374.0	1948.8	0.713521
KNN	2352.3	1917.1	31.7	21.7	2374.0	1948.8	0.990859

LSTM	1920.1	1515.5	433.3	453.9	2374.0	1948.8	0.808804
	TNR	FPR	FNR	Precision	F1		
Accuracy \							
Random Forest	0.989224	0.010776	0.009773	0.991148	0.990686		
0.989775							
Naive Bayes	0.538382	0.461618	0.286479	0.653139	0.681988		
0.634565							
KNN	0.983734	0.016266	0.009141	0.986706	0.988777		
0.987647							
LSTM	0.777657	0.222343	0.191196	0.816558	0.812118		
0.794762							
	Error Rate	Balanced Accuracy	TSS	HSS			
Brier Score \							
Random Forest	0.010225	0.989726	0.979452	0.979452			
0.010087							
Naive Bayes	0.365435	0.625952	0.251903	0.251903			
0.231035							
KNN	0.012353	0.987297	0.974593	0.974593			
0.009492							
LSTM	0.205238	0.793230	0.586461	0.586461			
0.141105							
	Brier Skill Score	ROC AUC					
Random Forest	0.959258	0.999267					
Naive Bayes	0.066833	0.687921					
KNN	0.961661	0.997119					
LSTM	0.430065	0.880433					
Comparison of Average Metrics Across Models:							
	Accuracy	F1	Precision	ROC AUC			
Random Forest	0.989775	0.990686	0.991148	0.999267			
Naive Bayes	0.634565	0.681988	0.653139	0.687921			
KNN	0.987647	0.988777	0.986706	0.997119			
LSTM	0.794762	0.812118	0.816558	0.880433			



## Comparison of Average Metrics Across All Algorithms

The table below summarizes the average performance metrics of the Random Forest, Naive Bayes, K-Nearest Neighbors (KNN), and LSTM models across all iterations:

Algorithm	T	T	F	F	T	T	P	P	Ac	Err	Balanc	T	H	Brie	Brier	R
	P	N	P	N	P	N	R	R	cu	or	ed	S	S	r	Skill	O
									ra	Rat	Accura	S	S	Score	Score	A
									cy	e	cy	S	S			U
																C
Random Forest	23	1	2	2	0	0	0.9	0	0.	0.0	0.9897	0	0	0.01	0.9593	0.
	5	9	1.	3.	.9	.9	911	.9	98	102		.9	.9	01		9
	0.	2	0	2	9	8		9	98			7	7			9
	8	7.			0	9		0				9	9			9
Naive Bayes		8			2	2		7				5	5			3
	16	1	8	6	0	0	0.6	0	0.	0.3	0.6260	0	0	0.23	0.0668	0.
	9	0	9	8	.7	.5	531	.6	63	654		.2	.2	10		6
	3.	4	9.	0	1	3		8	46			5	5			8
	9	9.	6	.1	3	8		2				1	1			7
		2			5	4		0				9	9			9

Algorithm	TP	TN	FP	FN	TPR	TNR	Precision	F1	Accuracy	Error Rate	Balanced Accuracy	TS	HS	Brier Score	Brier Skill Score	ROC AUC
<b>KNN</b>	2352.3	197.1	317	217	0.99	0.98	0.9867	0.987	0.99876	0.0124	0.9873	0.97	0.97	0.0095	0.9617	0.999
<b>LSTM</b>	1920.1	153.5	433	499	0.808	0.717	0.8166	0.81	0.7948	0.2052	0.7932	0.58	0.58	0.1719	0.5244	0.6879

## Detailed Comparison

### Random Forest: The Best Performer

- Performance:** The Random Forest model outperformed all other algorithms across all metrics.
  - It achieved the highest **Accuracy (98.98%)**, **Precision (99.11%)**, **F1-Score (99.07%)**, and **Balanced Accuracy (98.97%)**.
  - Its **True Positive Rate (TPR)** and **True Negative Rate (TNR)** were nearly perfect at 99.02% and 98.92%, respectively.
- Reliability:** The **Brier Score (0.0101)** indicates highly calibrated probability predictions, and the **Brier Skill Score (0.9593)** demonstrates its superior performance compared to a baseline model.
  - The **ROC AUC (0.9993)** confirms its ability to distinguish between classes effectively.
- Conclusion:** Random Forest is the most consistent and reliable algorithm for this task.

### Naive Bayes: The Weakest Performer

- Performance:** Naive Bayes struggled significantly in comparison to the other models.
  - It achieved the lowest **Accuracy (63.46%)**, **Precision (65.31%)**, and **F1-Score (68.20%)**.
  - Its **True Positive Rate (TPR)** was moderate at 71.35%, but the **True Negative Rate (TNR)** was very low at 53.84%, indicating poor handling of negative samples.
- Calibration:** The **Brier Score (0.2310)** and **Brier Skill Score (0.0668)** show that the probability estimates were poorly calibrated.
  - The **ROC AUC (0.6879)** indicates limited capability in distinguishing between classes.

- **Conclusion:** Naive Bayes, while simple and fast, is not suitable for this classification task due to its assumptions and inability to handle complex relationships in the data.

### K-Nearest Neighbors (KNN): A Close Second

- **Performance:** KNN performed exceptionally well and closely matched the Random Forest model:
  - It achieved an **Accuracy (98.76%)**, **Precision (98.67%)**, and **F1-Score (98.88%)**, with slightly lower scores than Random Forest.
  - Its **True Positive Rate (TPR)** and **True Negative Rate (TNR)** were 99.09% and 98.37%, respectively, showcasing high reliability.
- **Calibration:** The **Brier Score (0.0095)** and **Brier Skill Score (0.9617)** indicate that KNN produced well-calibrated probability predictions.
  - The **ROC AUC (0.9971)** shows excellent class discrimination.
- **Conclusion:** KNN is a strong competitor to Random Forest, with near-identical performance but slightly less consistency.

### LSTM: Moderate Performance

- **Performance:** The LSTM model showed moderate performance compared to the Random Forest and KNN models:
    - It achieved an **Accuracy (79.48%)**, **Precision (81.66%)**, and **F1-Score (81.21%)**.
    - Its **True Positive Rate (TPR)** was 80.88%, and the **True Negative Rate (TNR)** was 77.77%, reflecting reasonable performance but room for improvement.
  - **Calibration:** The **Brier Score (0.1719)** and **Brier Skill Score (0.5244)** indicate that the LSTM's probability predictions were less reliable compared to Random Forest and KNN.
    - The **ROC AUC (0.8865)** shows moderate class discrimination.
  - **Conclusion:** While LSTM leverages sequential data effectively, it falls short in terms of overall performance and consistency, potentially due to insufficient data or suboptimal hyperparameters.
- 

## Final Verdict

- **Best Performer: Random Forest**, with the highest accuracy, precision, and reliability.
- **Runner-Up: KNN**, with comparable performance to Random Forest but slightly lower calibration and consistency.
- **Moderate Performer: LSTM**, showing potential but requiring further optimization.
- **Weakest Performer: Naive Bayes**, struggling with accuracy, calibration, and handling of negative samples.

Based on this analysis, **Random Forest** is the most suitable algorithm for the sentiment analysis task in this project.

## Why Random Forest is the Best

Among the evaluated models, **Random Forest (RF)** emerges as the best performer for several reasons:

1. **Accuracy and Robustness:** Random Forest achieves the highest accuracy among the tested models, demonstrating its ability to handle the complexity and non-linear relationships in the mushroom dataset.
2. **Feature Importance:** It provides insights into feature importance, helping identify which mushroom characteristics are most predictive for classification.
3. **Resistance to Overfitting:** By aggregating the results of multiple decision trees, Random Forest reduces the risk of overfitting, even with complex datasets.
4. **Versatility:** It performs well across a wide range of data distributions and can handle both numerical and categorical features effectively.

The combination of high performance, interpretability, and robustness makes Random Forest the most reliable and practical choice for this binary classification task.

```
# Plotting ROC Curves for All Models
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Dictionary of models and their respective data
models = {
    'Random Forest': (rf_model, rf_metrics_df),
    'Naive Bayes': (nb_model, nb_metrics_df),
    'KNN': (knn_model, knn_metrics_df),
    'LSTM': (lstm_model, lstm_metrics_df) # Added LSTM
}

plt.figure(figsize=(10, 8))
for name, (model, metrics_df) in models.items():
    if name == 'LSTM':
        # Get predicted probabilities for LSTM
        y_prob = lstm_model.predict(X_test_lstm).flatten()
    else:
        # Get predicted probabilities for other models
        y_prob = model.predict_proba(X_test)[:, 1]

    # Compute ROC curve and AUC
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.2f})')

# Plot diagonal line for random guessing
plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
plt.title('ROC Curves for All Models')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.grid()
plt.show()
```



